

Forecasting the 2026 Sveriges Riksbank Prize in Economic Sciences

Record of the Nobel Prize/Econ pipeline through step 4a (the field forecast prompts)

Dawoon Jeong (with Claude Code)

10 October 2026

Abstract

This document records the work done in `/project/jevans/Dawoon/NobelPrize/Econ/` between 9 October 2026, 17:42 CDT, and 10 October 2026, about 11:30 CDT: the crawl of the complete prize record 1969–2025, the profiles of the eleven members of the 2026 prize committee written by three language models and merged, the JEL coding of the 64 awarded works and the analysis of fields, sequence, gender and age, and the construction of the Preseen context notes for a forecast of the *field* of the 2026 prize under a design that applies field rotation strongly. It states every design decision, the checks that were run, the numbers obtained, the costs, and what remains to be done before the announcement on Monday 12 October 2026 (11:45 CEST at the earliest, 04:45 CDT); forecasts must be finished by Sunday 11 October, 23:00 CDT. Later on 10 October it was extended with the results of the field forecast, the first reference facts for the people stage (age at the award, with a data correction) and the design of the Econ virtual committee that nominates candidates for the five leading fields.

Contents

1	Purpose, plan and status	2
2	Shared infrastructure	2
3	The prize record 1969–2025 (data/)	2
4	Step 1: profiles of the 2026 committee (01_committee/)	3
4.1	Roster	3
4.2	Design and checks	3
4.3	Results	3
5	Step 2: the fields of the prize (02_fields/)	4
5.1	Design	4
5.2	Results	4
6	Step 4a: the Preseen context notes for the field forecast (04_field_forecast/)	6
6.1	The design document	6
6.2	Reconciling the classifications	6
6.3	The record in the 14 fields	7
6.4	Committee coverage of the fields	7
6.5	The notes	8
6.6	The question	8
6.7	Treatments and arms	8
6.8	Execution	9
6.9	Results of the field forecast	9
7	Step 5: reference facts from the laureate record (05_laureates/)	10

8 Step 6: the Econ virtual committee (06_candidates/)	10
8.1 Purpose and design choice	10
8.2 Fields, prompt and schema	10
8.3 Validation, storage, aggregation, pool	11
8.4 Integration by Claude and the context notes	12
8.5 First run (five fields, 10 October 13:05–13:47 CDT; archived)	13
8.6 Second run (10 October, from 14:19 CDT)	13
9 Relation to Dolton and Tol (2026)	14
10 Open items and next steps	14
A Folder layout	14
B Commands	15
C Costs	15

1 Purpose, plan and status

The Econ folder is a forecast model for the economics prize separate from the persona experiments of `../experiment/`, which covered Physiology or Medicine, Physics and Chemistry with virtual committees whose members carried only the specialties of each committee and with options built around discoveries. The Econ model starts from the members of the real 2026 committee and from the complete award record, and proceeds in four steps: (1) profiles of the committee members, which may later become personas; (2) the fields the prize has gone to; (3) candidates per field; (4) the Preseen forecast. Step 4a, the subject of the last part of this record, builds the prompts for a Preseen question on the field; the people stage follows.

Timeline. Folder created 9 October 2026, 17:42 CDT. Crawl, step 1 and step 2 were built and run between 17:47 and 19:03 CDT the same day; the paid API calls ran 23:44–23:55 UTC (18:44–18:55 CDT). The step-4a notebook was written and executed on 10 October between about 11:00 and 11:30 CDT after the design document of 9 October was received.

2 Shared infrastructure

`config.yaml` holds the dates (information cutoff `today: 2026-10-09`, announcement, runs-finished-by), the names of the environment variables that hold the API keys (`COMMITTEE_ANTHROPIC_API_KEY`, `OPENAI_API_KEY`, `GEMINI_API_KEY`, `PRESEEN_API_KEY`; the keys themselves are never stored or printed), the three models of the 1 October virtual committees, unchanged (`claude-opus-5-5`, `gpt-5.5-2026-04-23`, `gemini-3.1-pro-preview`), the provider letters A, O, G used in the merged profiles, and the prices used for cost estimates. `llm_providers.py` makes JSON-mode calls to the three providers without tools or web search. Everything runs with the Curvature interpreter called by path (`/project/jevans/Dawoon/env/Curvature/bin/python`, with `LD_LIBRARY_PATH` set to its `lib`).

3 The prize record 1969–2025 (data/)

`crawl_econ_prizes.py` fetches three sources one second apart: the nobelprize.org list page in its *All Years* view (the default view shows only 2020–2025), and the Nobel Prize API 2.1 endpoints `nobelPrizes` and `laureates` for category `eco`. The page gives each year’s motivation groups (a year with two motivations was divided between two works) and the overall motivation where

one exists (2025); the API adds each laureate's share, affiliation at the prize, gender, birth and death dates. Every page laureate is matched to an API record and the two motivation texts are compared.

Results of the crawl of 9 October 2026. 57 prizes, 99 laureates (as the page states), 7 prizes divided between two works, 3 women laureates; every motivation identical in page and API. Outputs: `econ_prizes.md` (every prize by decade with laureates, shares, motivations, affiliations), `econ_prizes_by_year.csv`, `econ_prizes_laureates.csv` (28 columns per laureate), `crawl_report.json`; the raw page and API files in `raw/` with the date. Licence: API data CC0; the list page is © Nobel Prize Outreach, so the raw HTML stays local (`.gitignore`). One gap: the API record of Oliver E. Williamson has no death date (he died on 21 May 2020); the analysis notebook adds it and the crawled files stay as fetched.

4 Step 1: profiles of the 2026 committee (01_committee/)

4.1 Roster

The eleven members of the Committee for the Prize in Economic Sciences in Memory of Alfred Nobel 2026 were taken from the `nobelprize.org` committee page (fetched 9 October, saved in `source/`); the page gives no affiliations, which were checked against each person's university profile the same day (`roster.yaml`, with the link).

4.2 Design and checks

Each of the three models describes every member *from its own knowledge, without web search* (`prompts/ask_system.md`, `ask_user.md`; JSON schema; one call per member and model), returning an identity check (recognized, confidence), fields with JEL codes, recent interests (period, basis), up to eight representative works (title, year, venue, co-authors, contribution, confidence), a persona (summary, research lens, methods and evidence, what they value in contributions, questions they would raise in deliberation, committee experience) and knowledge limits. Claude (`claude-opus-5-5`, effort high) then merges the three answers into one profile (`merge_system.md`, `merge_user.md`), keeping for every field, interest and work the letters of the models that state it and the conflicts, and writing an integrated persona with a note where the models differ. Rules in the prompts: professional information only; accuracy over completeness (a model that does not recognize the person says so instead of guessing from the name); a work is listed only when the model is confident the title exists; the persona is grounded in the record and does not predict whom the member will support; the merge may use only what the three answers say.

Checks in code: every answer and merge validates against its schema (one retry on invalid JSON); support letters in the merge must belong to answers that exist; every merged work is matched by normalized title (similarity ≥ 0.85) against each model's own list, the rendered badge is this checked support, a work no model listed is marked, and a difference from the merge's own claim is shown. Every call is logged (stage, member, model, seconds, tokens, USD).

4.3 Results

All eleven members were recognized by all three models with high confidence; no namesake problems were found. The merged profiles are in `profiles/<slug>.md|json`, all on one page in `profiles/ALL.md`, one row per member in `profiles/index.csv`. 44 calls (33 answers, 11 merges), \$4.66, 23:44–23:55 UTC on 9 October.

Known limits: what a model knows stops at its training data (roughly 2024), so recent roles and papers can be missing or outdated; a single-model work needs checking. An earlier version of this step with web search (research, extract, aggregate) was built and tested offline but never

run, and was set aside in `Old/websearch_2026-10-09/` after the user chose the own-knowledge version.

5 Step 2: the fields of the prize (02_fields/)

5.1 Design

A *work* is one motivation group of a prize year: the 57 prizes, 7 of them divided, give 64 works. Each of the three models gives every work JEL codes from its own knowledge (`classify_jel.py`; 4 batches of 16 consecutive works, the works of a divided prize always in the same batch; the full JEL list from `data/classificationTree.xml` in the prompt: 20 letters, 139 level-2 and 856 level-3 codes). Per work the model returns a *primary* level-3 code (the core of the recognized contribution), up to three *secondary* codes, the kind of contribution (theory, theory and empirical, empirical, methods), whether it knows the prize beyond the motivation, and a one-sentence rationale. Rules: code the contribution the prize recognized, not the laureates' careers; level-3 codes from the list only; a tool that others apply (econometric, experimental, mathematical) goes under C with its first field as a secondary code. Checks in code: the JSON schema; every work id of the batch exactly once; every code a level-3 code of the tree, at most three secondary codes, distinct and not the primary; an invalid answer is asked once more with the problems listed. All 12 answers were valid at the first attempt (9 October, \$1.52).

Consensus (`prize_history.ipynb`, section 3), level by level (letter, then level 2 within it, then level 3 within that): the code that at least two of the three primary codes share; otherwise the highest score (2 per primary, 1 per model that lists it as secondary); a remaining tie goes to the first model in the order A, O, G.

5.2 Results

The three models agree on the letter for 58 of the 64 works (91%; pairwise kappa 0.90–0.96), on the level-2 field for 51 (80%) and on the level-3 code for 61%; every work has a two-model majority at the letter and level-2 levels. Six works divide the models at the letter level: 1984 Stone (E/C), 1997 Merton–Scholes (G/C), 2009 Ostrom (D/Q), 2010 Diamond–Mortensen–Pissarides (J/D), 2019 Banerjee–Duflo–Kremer (O/C), 2025 Mokyr (N/O).




Figure 1: Fields of the prize works by JEL letter (left) and by period (right); `02_fields/results/`.

Shifts. E was the main field eight times between 1974 and 2006 and not since (20 years against a median gap of 3), although macroeconomics is a secondary code of 2011 (Sargent, Sims) and 2022 (Bernanke, Diamond, Dybvig) for all three models. The last ten prizes: 2016 D, 2017 D, 2018 O+Q, 2019 O, 2020 D, 2021 C+J, 2022 G, 2023 J, 2024 O, 2025 N+O. Growth and development, labour and empirical work have gained; general equilibrium (D5, last 1988) and mathematical modelling (C5, C6, last 1980) have not returned.

Sequence. No sign that the committee avoids, or favours, the field of the year before: 10 years share a letter with the previous year against 10.2 expected in a random order of the 57 yearly sets ($P = 0.56$; the same for two- and three-year windows; 20,000 permutations). At level 2 a field followed itself only once (O4, 2024 $\rightarrow$ 2025; 1.5 expected). Theory and other work do not alternate more than chance (runs test $p = 0.11$), but empirical work rose to 3 of 8 works in 2020–25.




Figure 2: Left: the JEL letter of every prize work in award order (a half-size dot is half a prize). Right: time since the last award against the usual gap for letters with three or more award years.

Laureates, gender, age. The prize is shared more and more: 1.55 laureates per prize in 1969–79, 2.0 in the 2000s and 2010s, 2.5 in 2020–25 (one solo prize since 2020, Goldin 2023; the split $1/2 + 1/4 + 1/4$ only in 2021 and 2025). 82% of laureates had a US affiliation at the prize (at least 90% in every period since 2000); Chicago (15), MIT (10) and Harvard (8) lead. 3 women among 99 laureates: Ostrom (2009), Duflo (2019), Goldin (2023); since the first, 3 of 37 (8%). Median age at the award 67 (middle 80%: 56–77), no trend over time (+0.1 years per decade, $p = 0.81$); development and growth laureates are the youngest (median 63) and economic history and labour the oldest (72, 70), but the field differences are not significant (Kruskal–Wallis $p = 0.60$); laureates of three-way prizes are younger (median 64) than solo laureates (66.5) and those of two-way prizes (68). 48 of the 99 are alive (median age 79.4; 24 are 80 or older); the others died at a median 85.9, 18.2 years after their prize.

Data correction (10 October). The Nobel Prize API dates the 2022 prize (Bernanke, Diamond, Dybvig) 2011-10-10, so the crawl’s integer `age_at_award` was eleven years too low for these three laureates (57, 57, 56 instead of 68, 68, 67). The error was found when the step-5 notebook recomputed ages from the dates and checked them against the API values. The step-2 notebook now recomputes the age from the prize year where the API award date lies outside it (the crawled file stays as fetched, as for Williamson’s death date); the age figures above are the corrected ones (the first version of this record and the step-2 summary said median 66 and “finance youngest, median 57”, both artefacts of the error). The context note 00_5 uploaded to Preseen at 12:00 carries the API ages of the 2022 laureates; the ages are incidental to the field question. The uploaded texts are preserved in 04_field_forecast/uploaded_2026-10-10/ and the live notes are corrected.

The committee’s fields against the prize fields (section 9). A member covers a letter when one of their merged fields carries a code in it. Members covering / prizes 1969–2025: D 10 / 16.5, C 8 / 13, J 8 / 2.5, O 6 / 6, I 6 / 0, H 6 / 0, E 5 / 8, G 3 / 4, L 3 / 2, N 3 / 1.5, M 3 / 0, Q 2 / 0.5, A 2 / 0, K 2 / 0, R 2 / 0, F 1 / 3, B 1 / 0, P 1 / 0. The committee covers public economics and health, education and welfare, letters never awarded as the main field.

Outputs: `jel_by_model.csv`, `results/prize_works_jel.csv` (one row per work with the consensus at three levels, agreement, each model’s codes), `results/awarded_jel_fields.csv` |md, `results/field_gaps_letter.csv`, `field_gaps_level2.csv`, the figures, and the executed notebook (55 cells).

6 Step 4a: the Preseen context notes for the field forecast (04_field_forecast/)

6.1 The design document

On 10 October the user supplied a design document (written in Korean on 9 October, information date 9 October 2026, before the announcement): *Forecasting the field of the 2026 Nobel Prize in Economics: a design that applies field rotation strongly*. It marks every statement with an evidence tag — [P] a result reported by Dolton and Tol (2026), [M] an opinion raised at the project meeting of 9 October 2026 (transcript *E 60th St 25*, not in the repository), [D] a design choice, [O] official sources — and sets nine rules:

1. Field rotation is the most important prior: the record alone sets a baseline ranking, then the accumulation and maturity of unawarded work adjust it; societal attention and committee expertise are auxiliary; no field gets an automatic zero or a fixed turn. [D]
2. The same classification for the record and for this year's options: a project classification if one exists, otherwise the paper's 14 JEL-based fields; boundary rules fixed in advance; no re-labelling to avoid a penalty. [P, D]
3. Strong priority to the record with operational weights: awarded in the previous 1–2 years, large penalty; 3–5 years, penalty (more if concentrated); long wait, strong bonus; these tiers are the forecast's settings, not thresholds the paper estimated, and related facts are not double counted. [P, D]
4. The 2023–2025 awards (Goldin; Acemoglu, Johnson, Robinson; Mokyr, Aghion, Howitt) reflected exactly, with the AI/employment versus growth distinction and the institutions boundary handled consistently. [O, D]
5. Check that a long-waiting field has enough mature, unawarded work with living representatives (defining works, long programmes, textbooks, diffusion); the paper's candidate counts are not 2026 counts. [P, M, D]
6. Societal context (AI and jobs, trade and tariffs, war and migration, inequality, institutions and democracy) is auxiliary; topicality does not offset a recency penalty. [M, D]
7. Committee expertise and change over time are limited adjustments; roles are distinguished; no equal votes. [P, O, D]
8. The paper's numbers (coefficients, pseudo- R^2 , transition scores, individual-level regularities) are not this year's probabilities; information is limited to before the cutoff. [P, D]
9. A fixed output per field: representative contributions; last award, waiting time, recent frequency; baseline rank; maturity of unawarded work; the effect of context and committee; final rank and probability with grounds. [D]

The notebook renders the document in English and adds one tag, [S], for numbers computed from this project's data. The user asked for a notebook in which every cell prints the text of one Preseen context note, for the committee-profile-based field analysis to be included, and for the evidence to come from the data already completed.

6.2 Reconciling the classifications

Rule 2 keeps a project classification when one exists. Step 2 coded the record in JEL, which is finer than the paper's fields in some places (E has seven awarded level-2 fields) and coarser in others (D lumps general equilibrium, games, information and behavioural economics). The two are reconciled by taking the paper's 14 fields (Table B.7) as the options and mapping every work into them through its consensus JEL code:

Field	JEL codes (Table B.7; additions in italics)
Econometrics	C0–C5, <i>C8</i>
Equilibrium, Welfare	C6, D5, D6
Games, Market Structure	C7, D4
Behavioural, Experimental	C9, D9
Labour	D1, D3, I, J
Production, Industrial Organization	D2, L, <i>M</i>
Public, Law, Political Economy	D7, H, K, P
Information	D8
Macro	E
Trade	F
Finance	G
Development, Economic History	N, O1, O2, <i>O5</i>
Growth	O3, O4
Resources, Environment	Q, R

Decisions [D]: C8, M and O5, absent from Table B.7, were added so that every code has a field. One work’s consensus primary code is still outside the table: Ostrom 2009 (D02, institutions); it takes the field most often named by the three models’ secondary codes — Public, law and political economy (H41, D71: 5 of 7 mentions; Q20: 2). The 2024 prize (O43, institutions and growth) is Growth, with its secondary codes (P16, N40, D02) showing the overlap with Public and Development; mechanism design (D82) is Information, auction and matching design (D44, C78) Games; household, health and education economics are Labour; urban and regional economics are Resources. The notebook asserts that all 64 works map and that the prizes sum to 57.

6.3 The record in the 14 fields

Table 5 is the baseline from the record alone (rule 3): the fields ordered by years since the last award, ties broken by fewer prizes in 2016–2025 and then fewer prizes overall; the tier is the design’s operational weight. It differs from the JEL-letter view of step 2 in instructive ways: Microeconomics (D, 16.5 prizes, last 2020) splits into Equilibrium and welfare (last 1998), Games (2020), Information (2016), Behavioural (2017) and parts of Labour, Production and Public; the letter C splits into Econometrics (last 2021) and the mathematical methods of the 1970s, which join Equilibrium. The same field followed itself in 3 of 56 transitions (1973 Equilibrium, 1985 Macro, 2025 Growth), against 10 of 56 at the letter level.

6.4 Committee coverage of the fields

Each member’s profile fields carry JEL codes; a member covers a field when one of their profile fields has a code in it, and the mark is filled when at least two of the three models state that profile field. The main field of a member is the field with the largest support-weighted number of codes. For comparison, Dolton and Tol’s Table A.6 assigns one field and one JEL code to every committee member of 1969–2025, including the 2025 roster, which has the same eleven people as 2026 (Table 6).

Labour (9 members) and Public, law and political economy (8) are the most covered fields; Trade is covered by one member (Friberg, through pricing-to-market and exchange-rate pass-through); Econometrics, Equilibrium, Information, Macro, Development and Growth by five each. Main fields by the profiles: Hassler and Krusell Macro, Andersson Games, Dreber Almenberg Behavioural, Fredriksson and Hjalmarsson Labour, Strömberg Finance, Boppart Growth, Enflo Development and economic history, Friberg Production and IO, Teorell Public. The paper’s assignment differs for Boppart (Trade, F2, against growth and structural change in the profiles) and Friberg (Labour, I2, against industrial organization and international pricing); for Hassler and Krusell it says Growth (O4) where the profiles put macro first with growth, climate and public economics beside it. The notes use the coverage only as the limited adjustment of rule 7.

6.5 The notes

What was checked in the paper. The arXiv PDF (4.3 MB) was fetched on 10 October and its text extracted to read the figures quoted in the notes: Table B.7 (the 14 fields with candidates and laureates per field), Table 2 (field-level logit over 798 field-years: number of candidates 19.84, $p < 0.001$; transition probability with and without Lindbeck 75.76, $p < 0.001$; years since the field last won 0.0576, $p < 0.01$; year 0.0249, $p < 0.05$; won in the previous year 0.066 with the matrix and -1.78 without, both insignificant; log-likelihood -137 against -198 and pseudo- R^2 37% against 9.1% without the matrix), Table A.6 (committee members' fields) and the text on the sample split at 1997 (Table E.12: optimal age 71.4 then 70.8; committee proximity significant only in the second half; citation indicators significant only in the first). The design document's "39 years after the PhD" was not checked and is not quoted.

An external-signals note, built and dropped. A tenth note was built in the first version of the notebook from `Data/convergence_2026.csv`, which holds, for all four prizes, people with convergent external signals collected before the 2026 science prizes; its 58 economics rows carry 40 Nobel symposium dossier entries (2022 inequality, 2022 social network analysis, 2026 international economics), 12 entries from 2026 web prediction lists, 9 Kalshi prices (KXNOBELECON: Pakes 0.26, Athey 0.16, Blundell 0.08, Barro 0.07, Fehr, Mas-Colell and Dasgupta 0.06, Rabin 0.05, Zingales 0.02), 4 Clarivate 2026 citation laureates (Woodford, Athey, Varian, Winter), 8 high-predictive-value prizes (BBVA Frontiers 2024 and 2025, Nemmers 2024 and 2026, Clark Medals 2024–2026) and 2 Swedish proximity signals, each person given a field by the author from general knowledge. On 10 October the user decided to exclude it entirely (not even with the weaker `consider` treatment): information about named people is the channel the design wants to keep out of the field stage, and the field assignments were the author's. The cell was removed from the notebook and the run script refuses to start unless `context/` holds exactly the nine `00_` notes.

6.6 The question

`questions/fields14.json`: multiple choice, "Which field of economics will the 2026 Sveriges Riksbank Prize in Economic Sciences recognize?", 14 options of the form `<field> (JEL <codes>): <three or four keywords>` (for example "Information (JEL D8): asymmetric information, contract theory, mechanism design, decision making under uncertainty"), mutually exclusive and comprehensive, no Other. Three points were settled in discussion with the user on 10 October before submission. (i) Keywords were added to the labels so that the forecaster and the resolver read the same boundaries; the fine print says the keywords are illustrative and the JEL code groups define the fields. (ii) The resolution procedure is pre-committed: the primary code of the 2026 motivation is determined by the same three-model JEL coding and consensus rule that coded 1969–2025 and mapped by the code groups of the labels, so no judgment is exercised after the announcement. (iii) A prize divided between fields resolves by the overall motivation when the announcement gives one, otherwise by the larger share, then by the contribution named first; the earlier draft (larger share, then first named) would have resolved the 2025 prize to Development and economic history through Mokyr's first-named half although the overall motivation, "innovation-driven economic growth", is Growth. A primary code outside the 14 fields resolves by the majority of the secondary codes; if no field can be assigned the question is annulled. The instruction note states the same event so that the forecast and the resolution refer to one thing.

6.7 Treatments and arms

Preseen attaches a context note with a treatment. In this project's earlier runs `assume_true` behaved as a premise to apply (a note excluding one person from the Medicine forecast was applied; the Chemistry instruction notes moved the ranking to Spearman 0.14–0.39 against control), while `consider` behaved as one input among many (person cards attached only as

`consider` left the Medicine and Physics forecasts close to control). The nine notes are therefore attached as `assume_true`: the brief (00_1) only bites as a premise, the classification and the record (00_2–00_5) define the event, and 00_6–00_9 are the rest of the design; the committee note is the weakest fact but its own text restricts it to a limited adjustment. Two arms with identical private questions were submitted on 10 October (`run_field.sh`, tag `nobel26-econ-fields`, one rep each): *control* without context, which gives Preseen’s own view of the field, and *main* with the nine notes. Possible further arms, not submitted: all notes as `consider` (how much the strong-rotation premise drives the result), a second rep of each arm (run-to-run noise).

6.8 Execution

The notebook `field_context_prompts.ipynb` (31 cells) was generated from a script and executed with `nbconvert` on the login node (kernel `curvature`, about two minutes, no paid calls). The notebook asserts the input counts (64 works, 99 laureates, 11 valid profiles), the complete mapping, the prize total, and that every person of the convergence file has an assigned field; it writes the notes, the question, `results/*.csv` and the two LaTeX tables of this record. Nothing was uploaded to Preseen.

6.9 Results of the field forecast

Both runs completed on 10 October: control in 27 minutes (4 subforecasts, 357 sources consulted), main in 44 minutes (4 subforecasts, 465 sources). Table 8 sets the two distributions beside the record-only baseline.

- **The brief was followed.** The main write-up reports the three distributions the instruction asked for: a record-only distribution from its own operational formula, $(\text{wait} + 5)^{0.8} \times \text{tier} / (1 + 0.10 \times \text{ten-year shares})$ with tiers 1.25 / 1.00 / 0.60 / 0.40, then a maturity factor M per field (Production and IO 1.75 for structural empirical IO; Macro 1.60 for the New Keynesian programme; Trade 1.45 for heterogeneous-firm trade; Public 1.40; Equilibrium 0.85 and Growth 0.90 as the only discounts), then small auxiliary factors for societal relevance, committee expertise and dated recognition (none above 1.08). It restates that strong rotation is a modelling assumption and that the paper’s transition matrix is not used as probabilities. Its final ranking correlates 0.94 with the record-only baseline; the five bonus-tier fields hold 67.8 % (control 43.0 %), the 2021–2023 tier 7.6 % (25.6 %), the 2024–2025 fields 1.6 % (5.5 %).
- **Preseen’s own view (control)** leads with Production and IO 14.6 %, Macro 14.5, Econometrics 12.3, Labour 9.5. It built its own contribution-first coding of the 57 prizes, which differs from the step-2 coding in several assignments (Ostrom to Resources, the 2010 search prize to Information, the 2024 prize to Growth with Development at three awards), and it put the historical component below the contribution assessment.
- **Where the arms agree and differ.** Macro is first in both arms (16.8 against 14.5) and Production and IO is in the top three of both; the notes move Public (3.4 to 12.8), Trade (6.8 to 14.0) and Equilibrium (3.7 to 10.3) up and Econometrics (12.3 to 3.1) and Labour (9.5 to 2.3) down. Both distributions stay flat: no field exceeds 17 %.
- **External signals entered anyway.** Although the external-signals note was excluded, the main run found the same signals by its own search (Woodford’s Clarivate 2026 listing, the BBVA 2025 prize to Blanchard, Galí and Woodford, the Academy’s 2026 and 2027 Nobel symposia on international economics) and used them as the small factor R ; the control run used them more.
- **Resolution.** After the announcement on 12 October the motivation is coded with `classify_jel.py` (three models, consensus) and mapped by note 00_2; the log score of each arm over the 14 options is then recorded here. The write-ups are saved in `04_field_forecast/results/write_up_{control,main}.md`.

7 Step 5: reference facts from the laureate record (05_laureates/)

The people stage needs priors from the record of the 99 laureates. `laureate_history.ipynb` collects them section by section; each section ends with an English paragraph written to `prompt/<section>.md` for later use as a context note. Section 1 (10 October) covers the age at the award.

Age at the award. Age is computed from the birth date to the announcement (four laureates have a birth year only). Median 67 (bootstrap 95% 64–69), mean 67.3, standard deviation 8.3, middle 80% 57–78; 4 laureates under 55 (Duflo 47, Arrow 51, Merton 53, Kremer 55), 6 aged 80 or more (Hurwicz 90, Shapley 89, Schelling 84, Wilson 83, Vickrey 82, Coase 81). The distribution has one peak: a Gaussian mixture prefers one component by BIC, skewness +0.34, Shapiro–Wilk $p = 0.48$; the mode of the density is 65 (bootstrap 62–69). Two readings of “optimal age” are distinguished: the mode of the laureate distribution (65–67) and the age at which a *candidate* is most likely to win, which Dolton and Tol estimate on their candidate pool at 70–71 [P]. Without their pool the notebook computes the award rate among eventual laureates (of those not yet awarded at age a , the share awarded at a): it passes 5% at 61 and 10% at 68 and peaks around 77 among ages with at least ten laureates still unawarded, but it rises mechanically as the pool empties. The honest reading is a broad plateau from 60 to 77 rather than a sharp optimum. No trend over time (+0.13 years per decade, $p = 0.80$; 1969–1997 median 66, 1998–2025 median 67). By field the medians run from Information 61 and Growth 64 to Labour 70 and Games 74 (Kruskal–Wallis $p = 0.37$); three-way prizes go to younger people (median 64) than two-way prizes (69) and solo prizes (67) ($p = 0.07$). Co-laureates of one work differ in age by a median of 9 years; 4 of 26 shared works have a gap of 15 years or more (Hurwicz–Maskin–Myerson 34, Roth–Shapley 29, Mirrlees–Vickrey 22, Hicks–Arrow 17). Of the 51 deceased laureates the median died at 86, 18 years after the prize; 7 died within five years; of the 20 awarded at 75 or older, 13 have died, after a median of 11 years. The paragraph for the people stage (`prompt/01_age_at_award.md`) concludes that a candidate’s age matters mostly at the extremes (under 55 and 80 or over are rare) and that co-laureates are usually of one generation. The notebook is where the API’s 2022 award-date error was found (the exact ages were checked against the API’s integer ages and three disagreed by eleven years); see the data correction in the step-2 section.

8 Step 6: the Econ virtual committee (06_candidates/)

8.1 Purpose and design choice

The people stage starts from candidate pools per field, larger for the fields with the higher forecast probability. As for Physics, Chemistry and Medicine, a virtual committee nominates the candidates; unlike those committees, whose personas carried only the specialties of the real members, the Econ committee’s eleven personas are the real 2026 members as reconstructed in step 1. The user chose, on 10 October, the design in which one call per member and model covers all five fields (33 cells) over the alternative of one call per member, model and field (165 cells). The nominations of the three models are then integrated into one candidate list per field by `claude-opus-5-5` (the user’s choice over the rule-based clusters and a hand-made `merges.yaml`), and the lists become Preseen context notes for the people question. `virtual_committee.py` is written for the Econ pipeline and imports nothing from `experiment/`; the functions of the 1 October script that apply unchanged were copied into it (listed below).

8.2 Fields, prompt and schema

The five fields are those with the highest probability in the main (treatment) arm of the latest field run, the second experiment with the title “Which field of economics will the 2026 Nobel

Prize in Economic Sciences recognize?” (task 9b0a6220, finished 13:51 CDT): Macro 16.9%, Trade 15.6%, Production and IO 14.0%, Public, law and political economy 13.3%, Equilibrium and welfare 9.5% (`settings.yaml`; the first experiment had the same five in the same order, 16.8/14.0/13.9/12.8/10.3). On the user’s request the committee was briefly set up for all 14 fields in three batches (5/5/4 fields per call; the nine further definitions follow the boundary rules of note 00_2 and the awarded record), then restricted to the five leading fields; the other nine stay in `settings.yaml` with `batches_to_ask: [1]` and enter the prompt only as the list of the 14 fields with their codes, which fixes the boundaries. Each field carries a definition with its boundary rules (from note 00_2) and the list of the works already awarded in it, with the official motivations and the consensus code of step 2, from `04_field_forecast/results/prize_works_fields14.csv`.

The system prompt begins “You take the perspective of {name}, {role} of the Committee for the Prize in Economic Sciences in Memory of Alfred Nobel 2026 ({title}, {affiliation}). The perspective is reconstructed from the member’s public research record, not from any statement by the member; it is a simulation, and your nominations are inferences from that record, not the real person’s views.” It continues with the profile (summary, research lens, methods and evidence, what the member values, the questions the member would raise, fields with codes, recent interests) and ends: “Use this lens when you judge contributions. But, as the real committee does with the expert advisers it consults for areas outside its members’ own fields, nominate from the whole of each field, not only from the areas closest to this member’s work.” The user prompt gives the date (9 October 2026, before the announcement), the rules (a contribution of outstanding importance; at most three living people; no previous laureate; divided prizes possible; awarded contributions excluded and not to be re-labelled; the age record of step 5 as context on maturity, not as a filter), the list of the 14 fields with their codes, the five fields with definitions and awarded works, and the task: for each field up to five ranked nominations with `rank`, `contribution` (motivation style), `jel_code`, `people` (name, affiliation, born), `key_works` (≤ 3 defining publications, “the works that established the contribution”, only if known) and `rationale` (≤ 3 sentences from the member’s lens: why this contribution and these people, and the main reservation, if any), returned as one JSON object with the five field names exactly as given and the statement that an answer with fewer than five blocks is invalid. The prompt is 14.3k characters (`prompts/example_john-hassler_b1.md`).

The age paragraph replaced the first run’s single sentence (median 67, middle 80% 57–78, 4 of 99 under 55, 6 of 99 at 80 or more). It adds, from step 5: the yearly award rate among eventual laureates (under 5% before 61, over 10% from 68), the maximum of Dolton and Tol’s candidate-pool model at 70–71, no trend since 1969 and no reliable difference between fields, co-laureates of one generation (median gap 9 years; 15 or more in 4 of the 26 shared awards), 20 awards at 75 or older, and no posthumous award. Left out on purpose: the field medians (not significantly different, small n), the weak three-laureate effect ($p = 0.07$) and any preferred age band, which would act as a filter and be counted twice if Preseen also sees ages.

8.3 Validation, storage, aggregation, pool

Hard requirements: an object with a `fields` list; every one of the five fields present exactly once (matched by name, accent-insensitive) with at least one nomination; every nomination with the six keys and the right types. Warnings: ranks not 1.. n (re-ranked by stated rank, then position), more than five nominations (ranks above five earn no points), people outside 1–3, more than three works, a rationale of more than two sentences, a code outside the field’s groups. An invalid cell is asked once more and then recorded as missing. Each cell is stored once with the prompt, the schema, every attempt’s response metadata and usage, and the validated nominations by field (`committee/raw/<model>/<member>.json`); every call is logged with tokens and cost (`logs/calls.jsonl`).

`collect` splits every cell into one ballot per field (`committee/<field>/ballots.jsonl`). Aggregation per field is that of 1 October: rank r earns $6 - r$ points, each model’s points are

divided by its number of valid ballots in the field and summed over the models; nominations that share a person (first initial + surname, accents folded, split by differing middle initials) are linked into one option, and a reviewed `merges.yaml` can split, merge, alias, choose the wording or the shown people, and record deceased people. The option text is the wording of the nomination with the most normalized points; the people are ranked by normalized points; the code is the most frequent one. New in Econ: each option lists its supporters by member and role (chair, member, co-opted, secretary; no weights) and carries flags for a person who is already a laureate of this prize, a person who is a member of the 2026 committee (a sitting member is not awarded; such options are left out of the pool), a person born 1941 or earlier (age 85 or more in 2026), a contribution whose wording resembles an awarded motivation (word overlap ≥ 0.35) and a most frequent code outside the field. `pool` takes the top options per field, 7 for Macro, 6 each for Trade, Production and IO and Public, 5 for Equilibrium (30 in all, in proportion to the main-arm probabilities), and lists the people who appear in more than one field (`results/pool.csv`, `pool.md`).

Reuse. Copied unchanged from `experiment/preseen/committee.py`: `fold`, `name_tokens`, `person_key`, `middles`, `assign_identities`, `same_person`, `content_words`, `jaccard`, `make_clusters`, the scoring (`n_valid`, `score`), the shape of `describe`, `load_manual` and the `merges.yaml` conventions. New: the five-field ballot and its validation, the persona built from a profile, `collect`, the field definitions and awarded works from the Econ steps, the committee-member and age flags, `pool`.

8.4 Integration by Claude and the context notes

`integrate` sends, per field, every valid nomination of the three models to claude-opus-5-5 in one JSON-mode call (no tools; provider default effort; up to two re-asks that list the errors of the previous answer). The nominations carry short anonymous ids (`n001`, ...) and ballot labels (`b01`, ...; neither member nor model is shown), the rank on the ballot, the code, the contribution, the people with affiliation and stated birth year, the key works and the rationale; the people named in the nominations who cannot be awarded (sitting committee members, previous laureates) are listed. The system prompt asks Claude to merge the nominations into distinct candidates “as a committee secretary would before the deliberation”, using only what the nominations say: a candidate is one contribution that a single motivation could award to at most three people; nominations that describe the same body of work go together even if worded differently or naming a different subset of people; different contributions stay apart even if they share a person (no umbrella candidates); two nominations of the same ballot are joined only if they plainly describe the same contribution; every id goes into exactly one candidate. Per candidate Claude returns the ids, a motivation line, the JEL code, the lineup (1–3 people named in its nominations, ineligible people left out), up to four defining works copied from the nominations’ key works with the lineup people they belong to, the reasoning (≤ 3 sentences from the rationales, including the main reservation) and a grouping note. Claude does not rank or score.

The answer is checked locally (every id exactly once, no unknown id, the lineup named in the candidate’s nominations, defining works present if the nominations list any; a lineup person who cannot be awarded and a work not found among the nominations’ key works are warnings and flags). Then, from the assignment alone: the model-balanced Borda score (maximum 15), the models and members that nominated the candidate, for each lineup person the number of the candidate’s nominations that name them, for each defining work the number of nominations that list it, the other people named, and the flags of the rule-based aggregation (laureate, committee member, age ≥ 85 , resemblance to an awarded motivation). The result carries a stamp of its input ballots; `context` and `pool` refuse a stale integration. Outputs: `committee/<field>/integrated.json` and `integrated.md` (every candidate with its nominations, for review), `committee/integrated_raw/<field>.json` (prompt, attempts, id map).

`context` writes the Preseen notes in English: `context/06_00_committee_method.md` (source,

integration, how support is counted, caveats: a simulation, birth years as stated by the models, living status not checked) and one note per field, 06_01_macro.md to 06_05_equilibrium.md, with every candidate ranked by score: motivation and code, lineup with birth year, affiliation and support, score with models, members and nominations per model, other people named at least twice, defining works with “listed by”, the committee’s reasoning, and notes on ineligible or very old people. `pool -source claude` takes the top candidates per field (7/7/6/6/4 in proportion to the latest main arm). The rule-based `aggregate` remains for comparison (`pool -source rule`).

Independence and framing. The committee sees no Preseen output, no field probabilities, no external signals and no candidate profiles. Its nominations are a simulation reconstructed from the members’ published records and are kept private; they are not statements about the members’ views.

8.5 First run (five fields, 10 October 13:05–13:47 CDT; archived)

The test cell for the chair (13:05 CDT) was valid on gemini-3.1-pro-preview at the first attempt: 25 nominations, 90 seconds, \$0.16 (2.9k input, 12.6k output tokens of which 7.4k reasoning). Its ranks 1–3 per field: Macro — Kiyotaki and Moore (credit frictions), Krusell and Smith (heterogeneous-agent macro; Krusell is the committee’s secretary, which the aggregation now flags), Woodford (New Keynesian framework); Trade — Eaton and Kortum, Melitz, Obstfeld and Rogoff; Production and IO — Berry, Levinsohn and Pakes, Grossman and Moore (property rights), Bloom and Van Reenen; Public — Persson and Tabellini, Auerbach, Saez; Equilibrium — Dasgupta, Shoven and Whalley, Mas-Colell. The same cell on claude-opus-5-5 and gpt-5.5 failed with an authentication error (HTTP 401) before any tokens were spent; `check_keys.py` returns 401 for both keys while the Preseen and Gemini keys answer 200. The step-1 and step-2 calls of 9 October (18:44–18:55 CDT) had succeeded, so the two keys changed after that. The eleven Gemini cells were started at 13:09 CDT; the Anthropic and OpenAI cells ran from the user’s terminal with the renewed keys (13:28–13:43 CDT; the editor’s shell kept the old keys). 31 of 33 cells were valid: claude-opus-5-5 returned only the Macro block for Hassler and Boppart, twice each. The rule-based clusters were too coarse even after tightening the link rule (a shared person links two nominations only with two or more shared people or a content-word overlap of at least 0.15): the Trade leader absorbed 102 of 155 nominations (Melitz with Eaton–Kortum and the political economy of trade policy). That led to the integration by Claude. Cost \$8.14 over 61 calls including tests. The run is archived in 06_candidates/01d/five_fields_2026-10-10/ (README there) and replaced by the second run, which uses the revised prompt (age paragraph, defining works, reasoning with reservation).

8.6 Second run (10 October, from 14:19 CDT)

Started by the user from a terminal with valid keys: `run_all.sh` checks that the three model keys answer 200 and starts `run_all_steps.sh` detached (log `run_all.out`): the 33 calls with nine workers (models interleaved), a second pass that asks the calls still missing once more, `collect`, `integrate`, `context`, `pool`. Estimated \$8–9 for the calls and \$2–3 for the integration.

Result (done 14:29 CDT). All 33 calls were valid at the first attempt (25 nominations each, 165 per field; the second pass had nothing to ask), \$8.30. The integration was valid at the first attempt in every field (no warnings; every defining work is among the key works of the candidate’s nominations), \$2.33 (46–51k input and 10–16k output tokens per field). Total for the run \$10.63; all step-6 calls so far \$18.76.

Compared with the rule-based clusters of the first run, the integration separates the contributions that the clusters had chained: in Trade the leader holds 30 of 165 nominations (first run: 102 of 155), and Melitz, Eaton–Kortum, Grossman–Helpman, Obstfeld–Rogoff, Autor–Dorn–Hanson

and Antràs are separate candidates; in Macro, Kiyotaki–Moore and the narrative identification of policy effects (Romer, Romer, Ramey) are separate. A person can lead several candidates (Saez: sufficient statistics, top incomes, and with Chetty administrative data). Flags: Bewley 85, Sonnenschein 86, Anderson 86, Posner 87, Calabresi 94; three Equilibrium candidates resemble the 1972 general-equilibrium motivation by word overlap (distinct contributions in substance). The pool (30 options, 7/7/6/6/4) has no person in two fields. The five notes total about 127k characters (the nine field-stage notes 44k); the living status of the shown people is not yet checked.

9 Relation to Dolton and Tol (2026)

What is adopted: the 14 fields as the unit of comparison; years since the last award as the central descriptive variable; the idea that the number of mature candidates in a field matters (as a qualitative judgment, rule 5). What is deliberately not adopted: the coefficients, the fitted transition matrix and the pseudo- R^2 as probabilities (rule 8); the paper’s finding that the previous year’s field is not significant, which the design overrides with a deliberate strong-rotation assumption (rule 3, stated as such in the notes). Where the two differ in measurement: the paper counts laureates per field over a hand classification of citations (Information 13, Econometrics 13, Macro 12, Finance 11, Growth 9), this project counts share-weighted prizes over a model-coded classification (Macro 8.0, Equilibrium 7.0, Econometrics 6.5, Production 4.5, Labour 4.5, Information 4.0, Games 4.0, Finance 4.0, Growth 4.0); the paper puts three laureates in Public where this project has 1.5 prizes (Buchanan, Ostrom’s half).

10 Open items and next steps

- **Further arms** before Sunday 11 October, 23:00 CDT, if wanted: all notes as **consider** (how much the **assume_true** premise drives the result), a second rep of each arm (run-to-run noise). Questions **a5e12fca** (control) and **4d47d56b** (main); tasks **09b31a2a** and **5d14390d**.
- **Check the English rendering** of the Korean design document in the notes against the original.
- **Length.** **00_3** (9.7k characters) and **00_8** (8.5k) are longer than the earlier notes of this project; the API accepted them.
- **After the announcement:** code the 2026 motivation as step 2 did, resolve the question, score the arms (log score over 14 options; the earlier runs scored Medicine at -2.54 to -2.90 and Physics at -3.23 to -3.86 with 12 options plus Other, uniform -2.56).
- **People stage** (steps 3 and 4b): candidates per field, the profiles as personas (1b) if the discussion design is wanted.
- **Housekeeping:** the folder is untracked in the Nobel Prize repository; **record/build/** is to be ignored; the **[M]** transcript is not on disk.

A Folder layout

```
Econ/
  README.md, config.yaml, llm_providers.py, .gitignore
  data/
    crawl_econ_prizes.py, econ_prizes.md, econ_prizes_by_year.csv,
    econ_prizes_laureates.csv, crawl_report.json, classificationTree.xml,
    raw/ (local)
  01_committee/
    roster.yaml, profiles.py, prompts/, settings.yaml,
    profiles/ (<slug>.md|json, ALL.md, index.csv),
    runs/ and logs/ (local), source/, Old/websearch_2026-10-09/
  02_fields/
    classify_jel.py, prompts/, settings.yaml, jel_by_model.csv,
    prize_history.ipynb, results/ (prize_works_jel.csv,
    awarded_jel_fields.*, field_gaps*.csv, figures), runs/ and logs/ (local)
  04_field_forecast/ README.md, field_context_prompts.ipynb, context/ (00_1..00_9),
```

	questions/fields14.json, run_field.sh, nobel_preseen_exp.py (client copy),
	results/ (fields14_history.csv, prize_works_fields14.csv,
	committee_fields14_*.csv), preseen_exp/ and *.out (local)
05_laureates/	README.md, laureate_history.ipynb, results/ (figures, age_*.csv), prompt/01_age_at_award.md
06_candidates/	README.md, virtual_committee.py, settings.yaml, prompts/, committee/raw/ (local),
	committee/test/ (local), committee/<field>/ (ballots.jsonl, candidates.json, review.md,
	merges.yaml), results/ (pool.csv, pool.md), logs/ (local)
record/	Record.tex (this document), tables/ (written by the step-4a notebook),
	build/ (local)

B Commands

```

cd "/project/jevans/Dawoon/Nobel Prize/Econ"
export LD_LIBRARY_PATH=/project/jevans/Dawoon/env/Curvature/lib
PY=/project/jevans/Dawoon/env/Curvature/bin/python
$PY data/crawl_econ_prizes.py                # prize record (free)
$PY 01_committee/profiles.py run              # step 1 (paid; valid records kept)
$PY 02_fields/classify_jel.py run && $PY 02_fields/classify_jel.py export
NB="$PY -m jupyter nbconvert --to notebook --execute --inplace"
NB="$NB --ExecutePreprocessor.kernel_name=curvature"
$NB 02_fields/prize_history.ipynb
$NB 04_field_forecast/field_context_prompts.ipynb
cd record && export PATH=/software/texlive-2023/bin/x86_64-linux:$PATH
export TEXMFVAR=/scratch/midway3/jdwoon0523/texmf-var2023b && mkdir -p build
for i in 1 2 3; do
  pdflatex -interaction=nonstopmode -output-directory build Record.tex > build/run$i.log
done

```

C Costs

Step	Calls	Count	USD	When (UTC, 9 Oct 2026)
1 ask	11 members $\times$ 3 models	33	4.66	23:44–23:54
1 merge	11 members, claude-opus-5-5	11		
2 JEL codes	4 batches $\times$ 3 models	12	1.52	23:49–23:51
4a notes	none (local computation)	0	0	10 Oct, 16:18–17:00
4b Preseen	2 forecast runs (control, main)	2	Preseen credits	10 Oct, 17:00–17:44
5 laureate facts	none (local computation)	0		10 Oct, 17:30–18:00
6 committee test	1 member $\times$ 3 models (2 auth failures)	3	0.16	10 Oct, 18:05–18:07
6 committee run	11 members $\times$ 3 models	33	about 8 (est.)	Gemini from 18:09; others pending
Total so far		59	6.34	

Step	Folder	Status on 10 October 2026
Prize record 1969–2025	data/	done 9 Oct: 57 prizes, 99 laureates, 64 works; list page and API agree on every year, name and motivation
1. Committee profiles	01_committee/	done 9 Oct: 11 of 11 merged; 44 calls, \$4.66
1b. Personas, discussion	–	not started; the field question uses the profiles as context notes instead (step 4a)
2. Fields	02_fields/	done 9 Oct: JEL codes of 64 works by 3 models (12 calls, \$1.52), consensus, analysis notebook executed to do (people stage)
3. Candidates per field	–	done 10 Oct: notebook executed, 9 notes, question with 14 keyword-labelled fields
4a. Context notes for the field question	04_field_forecast/	done 10 Oct 12:00–12:44 CDT: control (no context) and main (9 notes, <code>assume_true</code>), one rep each; main: Macro 16.8 %, Trade 14.0, Production/IO 13.9, Public 12.8, Equilibrium 10.3
4b. Preseen run	04_field_forecast/run_field.sh	section 1 (age at the award) done 10 Oct; found and corrected the API's 2022 award-date error
5. Laureate reference facts	05_laureates/	built 10 Oct: 11 real-member personas $\times$ 3 models, one call per cell over the five leading fields; Gemini cells running, Anthropic and OpenAI cells wait for valid keys
6. Virtual committee, candidates	06_candidates/	

Table 1: Status of the Econ pipeline. The folder is not yet tracked in the Nobel Prize git repository.

Member	Role	Affiliation
John Hassler	chair	IIES, Stockholm University
Tommy Andersson	member	Department of Economics, Lund University
Anna Dreber Almenberg	member	Department of Economics, Stockholm School of Economics
Peter Fredriksson	member	Department of Economics, Uppsala University
Per Strömberg	member	Stockholm School of Economics (Swedish House of Finance)
Timo Boppart	co-opted member	IIES, Stockholm University
Kerstin Enflo	co-opted member	Department of Economic History, Lund University
Richard Friberg	co-opted member	Department of Economics, Stockholm School of Economics
Randi Hjalmarsson	co-opted member	Department of Economics, University of Gothenburg
Jan Teorell	co-opted member	Department of Political Science, Stockholm University
Per Krusell	secretary	IIES, Stockholm University

Table 2: The 2026 committee (`roster.yaml`).

Member	Fields	Interests	Works	Listed by all three	Listed by one model
John Hassler	6	6	9	2	3
Tommy Andersson	7	6	11	0	8
Anna Dreber Almenberg	7	6	10	3	6
Peter Fredriksson	5	6	9	2	3
Per Strömberg	6	8	10	3	3
Timo Boppart	7	6	5	1	2
Kerstin Enflo	6	7	6	2	3
Richard Friberg	5	5	11	1	7
Randi Hjalmarsson	6	7	14	0	8
Jan Teorell	6	6	11	2	5
Per Krusell	8	5	13	3	10

Table 3: Merged profiles (`profiles/index.csv`): merged fields, recent interests and representative works per member, and how many works all three models or only one model listed.

Letter	Field	Works	Prizes (share-weighted)	Years
D	Microeconomics	18	16.5	1972 ... 2020
C	Mathematical and Quantitative Methods	16	13	1969 ... 2021
E	Macroeconomics and Monetary Economics	8	8	1974 ... 2006
O	Development, Innovation, Technological Change, Growth	7	6	1971, 1979, 1987, 2018, 2019, 2024, 2025
G	Financial Economics	4	4	1990, 1997, 2013, 2022
F	International Economics	3	3	1977, 1999, 2008
J	Labor and Demographic Economics	3	2.5	2010, 2021, 2023
L	Industrial Organization	2	2	1982, 2014
N	Economic History	2	1.5	1993, 2025
Q	Agricultural, Natural Resource and Environmental	1	0.5	2018

Table 4: Consensus JEL letter of the 64 works. Never the main field: A, B, H, I, K, M, P, R, Y, Z (H, I, K and P appear among the secondary codes). 35 level-2 fields have won; the most frequent are O4 growth (5 works), D8 information (4) and C2 single-equation econometrics (4).

Rank	Field (JEL)	Prizes	Last	Since	Median gap	2021–25	2016–25	Tier
1	Equilibrium, Welfare (C6, D5, D6)	7.0	1998	28	3.5	0.0	0.0	bonus
2	Macro (E)	8.0	2006	20	3	0.0	0.0	bonus
3	Trade (F)	3.0	2008	18	15.5	0.0	0.0	bonus
4	Public, Law, Political Economy (D7, H, K, P)	1.5	2009	17	23	0.0	0.0	bonus
5	Production, Industrial Organization (D2, L, M)	4.5	2014	12	7	0.0	0.0	bonus
6	Information (D8)	4.0	2016	10	6	0.0	1.0	neutral
7	Behavioural, Experimental (C9, D9)	2.0	2017	9	15	0.0	1.0	neutral
8	Resources, Environment (Q, R)	0.5	2018	8	–	0.0	0.5	neutral
9	Games, Market Structure (C7, D4)	4.0	2020	6	8	0.0	1.0	neutral
10	Econometrics (C0–C5, C8)	6.5	2021	5	9.5	0.5	0.5	penalty
11	Finance (G)	4.0	2022	4	9	1.0	1.0	penalty
12	Labour (D1, D3, I, J)	4.5	2023	3	5.5	1.5	1.5	penalty
13	Development, Economic History (N, O1, O2, O5)	3.5	2025	1	14	0.5	1.5	large penalty
14	Growth (O3, O4)	4.0	2025	1	11	1.5	2.0	large penalty

Table 5: The 64 prize works in the 14 fields (share-weighted prizes; years since the last award to 2026; median gap between award years; prizes in 2021–25 and 2016–25; operational tier of the design). Written by the step-4a notebook to `tables/fields14_history.tex`.

Member	Role	Em	Eq	Gm	Bh	Lb	IO	Pb	In	Ma	Tr	Fi	Dv	Gr	Re	Dolton–Tol	A.6
John Hassler	chair (member)		•			•		•		•				•	•	Growth (O4)	
Tommy Andersson	member		•	•		•		•	•						•	Information (D8)	
Anna Dreber Almenberg	member	•	◦	◦	•	•		•	•			•				Behavioural (D9)	
Peter Fredriksson	member	•				•	◦	•	•	•						Labour (I2)	
Per Strömberg	member			•			•	•	•			•	◦			Finance (G2)	
Timo Boppert	co-opted		•			•	•			•			•	•		Trade (F2)	
Kerstin Enflo	co-opted	•				◦				•			•	•	•	Development (N1)	
Richard Friberg	co-opted					◦	•	◦	•		•	•				Labour (I2)	
Randi Hjalmarsson	co-opted	•			•	•		•					◦			Labour (I2)	
Jan Teorell	co-opted	•						•					•	•		Public (D7)	
Per Krusell	secretary		•		◦	•		•		•				•	•	Growth (O4)	
members covering		5	5	3	3	9	4	8	5	5	1	3	5	5	4		

Table 6: Coverage of the 14 fields by the 2026 committee according to the step-1 profiles (• stated by 2–3 models, ◦ by one), with Dolton and Tol’s one-field assignment. Em econometrics, Eq equilibrium, Gm games, Bh behavioural, Lb labour, IO production and IO, Pb public, In information, Ma macro, Tr trade, Fi finance, Dv development and economic history, Gr growth, Re resources.

File	Content	Tags
00_1_instruction	goal, evidence tags, the rotation prior, the operational weights, pointers to the reference notes, the required output; written for the <code>assume_true</code> treatment	P M D
00_2_field_classification	the 14 fields and their codes, the additions, the boundary rules	P D
00_3_prize_works_by_field	every prize work 1969–2025 in its field: year, laureates, share, consensus code and JEL name, fallback rule, other fields the secondary codes point to	S O
00_4_prize_history_by_field	Table 5 as a markdown table; sequence facts	S
00_5_recent_prizes_2021_2025	the seven works of 2021–2025 with laureates, shares, affiliations, ages, official motivations, the three models’ codes, field and weight; how to read them	O S D
00_6_unawarded_maturity	what to assess; the paper’s pool sizes per field (Table B.7) and why they are not 2026 counts; what the project does not provide	P M D
00_7_societal_context	the meeting’s themes as auxiliary information	M D
00_8_committee_2026	roster and roles, the coverage table, one line per member, how to use it	O S P D
00_9_paper_evidence_and_limits	the paper’s field-level results and the limits of using them	P D

Table 7: The context notes in `04_field_forecast/context/`. Sizes 1.1k–9.7k characters, 44k in all; attached in name order, all with `treatment=assume_true`.

Field	Baseline rank	Tier	Control %	Main %	Main – control (pp)	Main rank
Macro	2	bonus:	14.5	16.8	+2.3	1
Trade	3	bonus:	6.8	14.0	+7.2	2
Production, Industrial Organization	5	bonus:	14.6	13.9	-0.7	3
Public, Law, Political Economy	4	bonus:	3.4	12.8	+9.4	4
Equilibrium, Welfare	1	bonus:	3.7	10.3	+6.6	5
Information	6	neutral	8.2	6.5	-1.7	6
Resources, Environment	8	neutral	3.8	6.3	+2.5	7
Behavioural, Experimental	7	neutral	6.3	5.8	-0.5	8
Games, Market Structure	9	neutral	7.5	4.4	-3.1	9
Econometrics	10	penalty	12.3	3.1	-9.2	10
Labour	12	penalty	9.5	2.3	-7.2	11
Finance	11	penalty	3.8	2.2	-1.6	12
Development, Economic History	13	large	2.7	0.9	-1.8	13
Growth	14	large	2.8	0.7	-2.1	14
bonus tier (5 fields)			43.0	67.8		
penalty tiers (5 fields)			31.1	9.2		

Table 8: Preseen’s probability per field, control (no context) against main (nine notes as `assume_true`), with the record-only baseline rank and tier of Table 5. Spearman main–control 0.43, main–baseline 0.94, control–baseline 0.27; entropy main 2.36, control 2.48, uniform 2.64.

	1 October committees (Med, Phys, Chem)	Econ
Personas	specialties of the real committee, anonymous	the 11 real members, from the merged step-1 profiles (summary, research lens, methods and evidence, what they value, questions they would raise, fields, recent interests)
Cells	persona $\times$ model, one ballot per field (66 cells)	member $\times$ model, one call over the five fields (33 cells)
Nominations	≤ 5 ranked per ballot	≤ 5 ranked per field, ≤ 25 per cell
Nomination	discovery line, 1–3 living people with affiliation, ≤ 3 key papers, rationale ≤ 2 sentences	the same, plus the JEL level-3 code of the contribution and each person’s year of birth
Models	claude-opus-5-5, gpt-5.5-2026-04-23, gemini-3.1-pro-preview; no tools, no web search, provider defaults, 32k output cap	the same
Validation	JSON schema, one retry, count limits as warnings	the same, plus every field present once, and a warning when a code lies outside the field’s groups
Aggregation	model-balanced Borda 5..1, person-key clustering, reviewed <code>merges.yaml</code>	model-balanced Borda 5..1 per field; grouping of the nominations into candidates by claude-opus-5-5 (<code>integrate</code>); supporters by member and role, per-person support and flags computed locally
Output	top- K options and “Other”	one Preseen context note per field (candidates with lineup, defining works, reasoning and support), and a 30-option pool allocated to the fields in proportion to the main-arm probability

Table 9: The Econ virtual committee against the 1 October committees.

Field	cand.	≥ 2 members	≥ 3 members	top-8 share	leading candidates (score of 15; members of 11)
Macro	26	18	11	0.81	New Keynesian framework: Woodford, Galí, Gertler (12.1; 11); credit constraints and collateral: Kiyotaki, Moore (8.5; 11); heterogeneous-agent macro: Huggett, Smith, Bewley (4.8; 10; Krusell, named most often, left out as committee member); Taylor rules and staggered contracts (3.2; 10)
Trade	18	11	10	0.96	firm heterogeneity: Melitz, Bernard, Helpman (11.5; 11); quantitative Ricardian models: Eaton, Kortum (7.3; 11); innovation, growth and the political economy of trade policy: Grossman, Helpman (7.1; 11); new open-economy macro: Obstfeld, Rogoff (4.9; 11)
Production, IO	28	18	13	0.81	demand and market power in differentiated products: Berry, Levinsohn, Pakes (14.5; 11); management practices: Bloom, Van Reenen, Sadun (7.5; 11); sunk costs and market structure: Sutton (4.8; 8); two-sided markets: Rochet, Armstrong, Jullien (4.2; 10)
Public, law, pol. econ.	22	15	13	0.82	political institutions and policy: Persson, Tabellini (10.4; 11); sufficient statistics for tax and social insurance: Saez, Chetty, Kleven (5.5; 8); economic analysis of law: Shavell, Posner, Calabresi (5.4; 11); top incomes and wealth: Saez, Piketty, Zucman (4.6; 8)
Equilibrium, welfare	24	16	12	0.81	structure and limits of general equilibrium: Mas-Colell, Sonnenschein, Balasko (8.5; 11); equality of opportunity and fairness: Roemer, Fleurbaey, Maniquet (6.9; 11); incomplete-markets equilibrium: Geanakoplos, Polemarchakis, Magill (6.0; 10); sustainability and natural capital: Dasgupta, Heal (4.7; 10)

Table 10: Integrated candidates of the second run. “top-8 share” = share of the field’s total score held by its eight leading candidates.